

# Assignment 5: Word Relationships

Desik Somasundaram

```
library(tidyr) #text analysis in R
library(pdftools)
library(lubridate) #working with date data
library(tidyverse)
library(tidytext)
library(readr)
library(quanteda)
library(readtext) #quanteda subpackage for reading pdf
library(quanteda.textstats)
library(quanteda.textplots)
library(ggplot2)
library(forcats)
library(stringr)
library(quanteda.textplots)
library(widyr) # pairwise correlations
library(igraph) #network plots
library(ggraph)
library(here)
```

## Import EPA EJ Data

```
setwd("dat/")
files <- list.files(pattern = "*.pdf$")

files <- str_subset(files, pattern="EPA")

ej_reports <- lapply(files, pdf_text)

ej_pdf <- readtext(files, docvarsfrom = "filenames",
                   docvarnames = c("type", "subj", "year"),
                   sep = "_")

#creating an initial corpus containing our data
epa_corp <- corpus(x = ej_pdf, text_field = "text" )
summary(epa_corp)
```

## Corpus consisting of 6 documents, showing 6 documents:

##

##		Text	Types	Tokens	Sentences	type	subj	year
##	EPA_EJ_2015.pdf	2136	8944	263	EPA	EJ	2015	
##	EPA_EJ_2016.pdf	1599	7965	176	EPA	EJ	2016	

```
## EPA_EJ_2017.pdf 3973 30564      653 EPA EJ 2017
## EPA_EJ_2018.pdf 2774 16658      447 EPA EJ 2018
## EPA_EJ_2019.pdf 3773 22648      672 EPA EJ 2019
## EPA_EJ_2020.pdf 4493 30523      987 EPA EJ 2020
```

```
#I'm adding some additional, context-specific stop words to stop word lexicon
more_stops <-c("2015", "2016", "2017", "2018", "2019", "2020", "www.epa.gov", "https")
add_stops<- tibble(word = c(stop_words$word, more_stops))
stop_vec <- as_vector(add_stops)
```

## Tidying and Exploring Data

Now we'll create some different data objects that will set us up for the subsequent analyses

```
#convert to tidy format and apply my stop words
raw_text <- tidy(epa_corp)

#Distribution of most frequent words across documents
raw_words <- raw_text %>%
  mutate(year = as.factor(year)) %>%
  unnest_tokens(word, text) %>%
  anti_join(add_stops, by = 'word') %>%
  count(year, word, sort = TRUE)

#number of total words by document
total_words <- raw_words %>%
  group_by(year) %>%
  summarize(total = sum(n))

report_words <- left_join(raw_words, total_words)
```

```
## Joining, by = "year"
```

```
par_tokens <- unnest_tokens(raw_text, output = paragraphs, input = text, token = "paragraphs")

par_tokens <- par_tokens %>%
  mutate(par_id = 1:n())


par_words <- unnest_tokens(par_tokens, output = word, input = paragraphs, token = "words")
```

Let's see which words tend to occur close together in the text. This is a way to leverage word relationships (in this case, co-occurrence in a single paragraph) to give us some understanding of the things discussed in the documents.

```
word_pairs <- par_words %>%
  pairwise_count(word, par_id, sort = TRUE, upper = FALSE) %>%
  anti_join(add_stops, by = c("item1" = "word")) %>%
  anti_join(add_stops, by = c("item2" = "word"))
```

Now we can visualize

```
word_pairs %>%
  filter(n >= 70) %>%
  graph_from_data_frame() %>%
  ggraph(layout = "fr") +
  geom_edge_link(aes(edge_alpha = n, edge_width = n), edge_colour = "dodgerblue") +
  geom_node_point(size = 5) +
  geom_node_text(aes(label = name), repel = TRUE,
    point.padding = unit(0.2, "lines")) +
  theme_void()
```



Hmm, interesting, but maybe we further subset the word pairs to get a cleaner picture of the most common ones by raising the cutoff for number of occurrences (n).

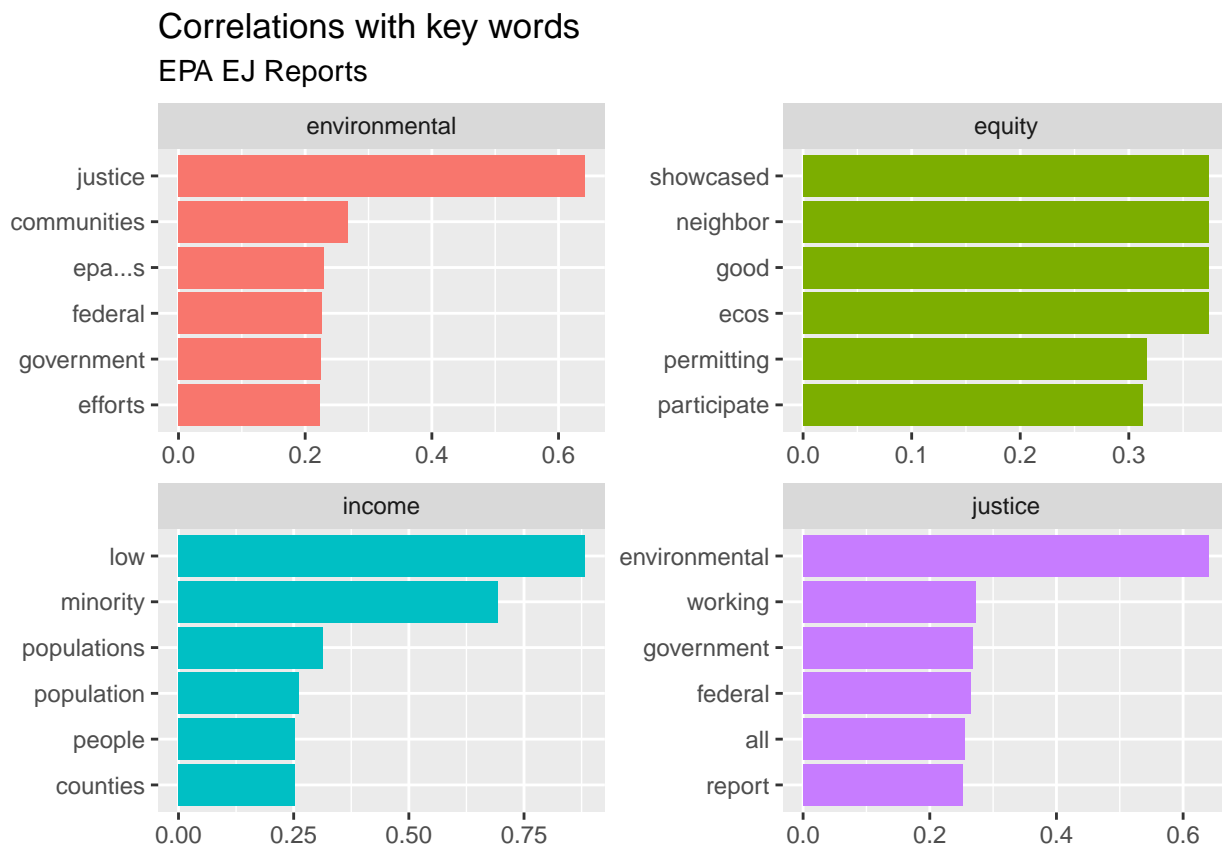
Pairs like “environmental” and “justice” are the most common co-occurring words, but that doesn’t give us the full picture as they’re also the most common individual words. We can also look at correlation among words, which tells us how often they appear together relative to how often they appear separately.

```
word_cors <- par_words %>%
  add_count(par_id) %>%
  filter(n >= 50) %>%
  select(-n) %>%
  pairwise_cor(word, par_id, sort = TRUE)

just_cors <- word_cors %>%
  filter(item1 == "justice")
```

```
word_cors %>%
  filter(item1 %in% c("environmental", "justice", "equity", "income"))%>%
  group_by(item1) %>%
  top_n(6) %>%
  ungroup() %>%
  mutate(item1 = as.factor(item1),
         name = reorder_within(item2, correlation, item1)) %>%
  ggplot(aes(y = name, x = correlation, fill = item1)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~item1, ncol = 2, scales = "free")+
  scale_y_reordered() +
  labs(y = NULL,
       x = NULL,
       title = "Correlations with key words",
       subtitle = "EPA EJ Reports")
```

## Selecting by correlation

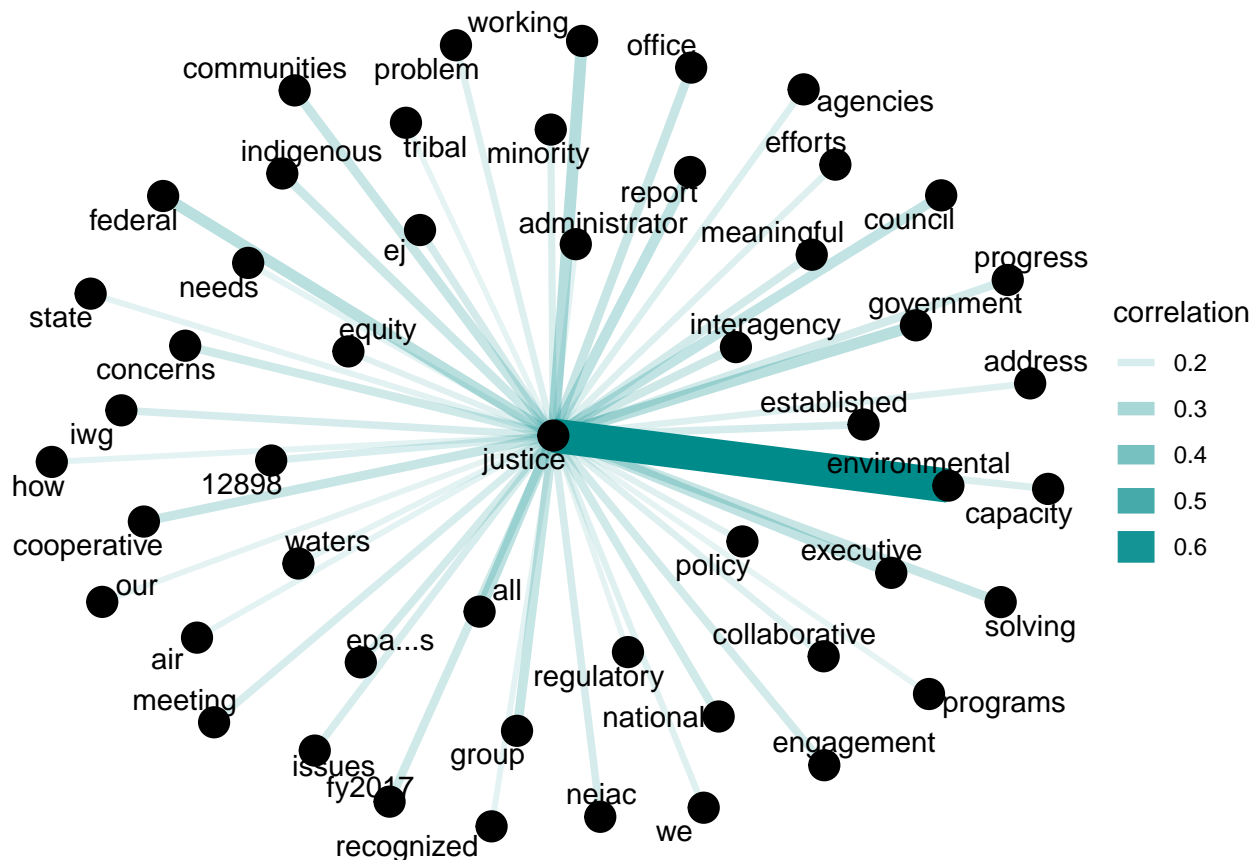


```
#let's zoom in on just one of our key terms
justice_cors <- word_cors %>%
  filter(item1 == "justice") %>%
  mutate(n = 1:n())
```

Not surprisingly, the correlation between “environmental” and “justice” is by far the highest, which makes

sense given the nature of these reports. How might we visualize these correlations to develop of sense of the context in which justice is discussed here?

```
justice_cors %>%
  filter(n <= 50) %>%
  graph_from_data_frame() %>%
  ggraph(layout = "fr") +
  geom_edge_link(aes(edge_alpha = correlation, edge_width = correlation), edge_colour = "cyan4") +
  geom_node_point(size = 5) +
  geom_node_text(aes(label = name), repel = TRUE,
                 point.padding = unit(0.2, "lines")) +
  theme_void()
```

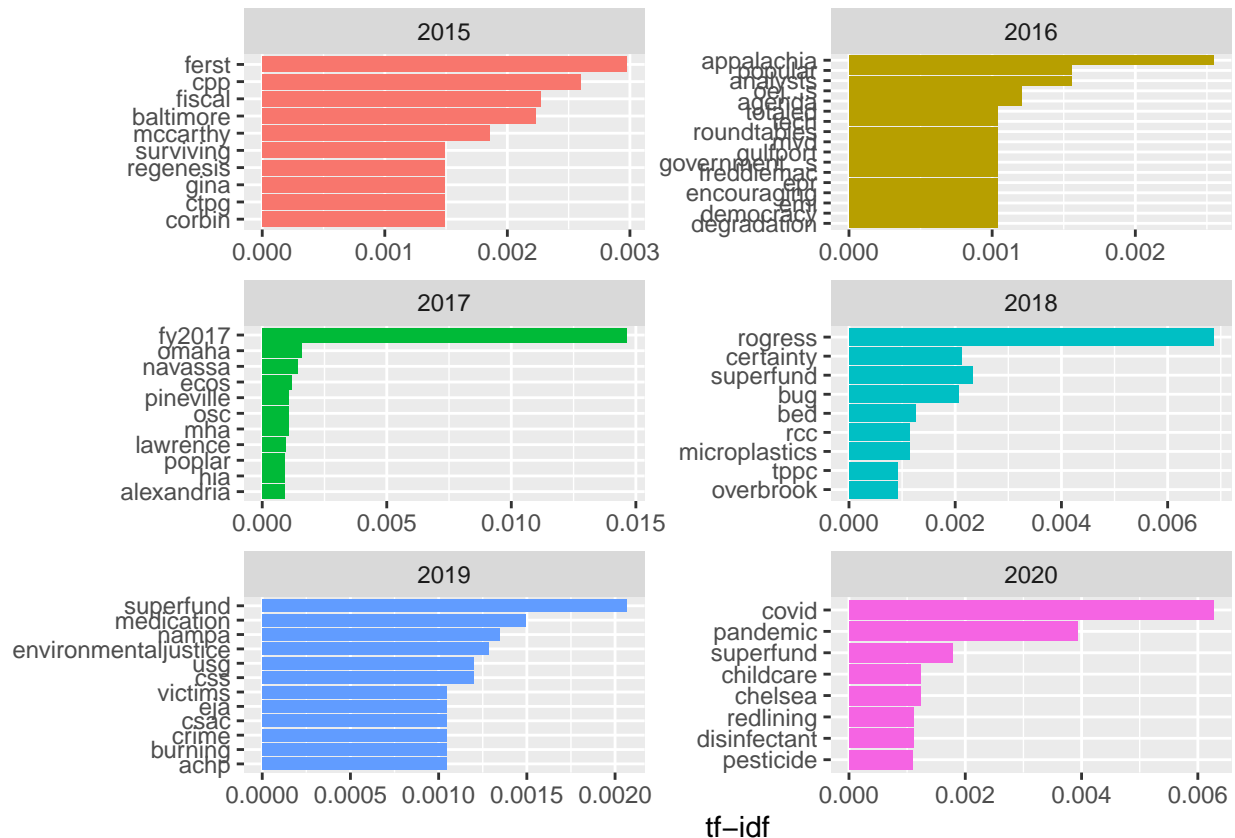


Now let's look at the tf-idf term we talked about. Remember, this statistic goes beyond simple frequency calculations within a document to control for overall commonality across documents

```
report_tf_idf <- report_words %>%
  bind_tf_idf(word, year, n) %>%
  select(-total) %>%
  arrange(desc(tf_idf))

report_tf_idf %>%
  group_by(year) %>%
  slice_max(tf_idf, n = 10) %>%
  ungroup() %>%
  filter(nchar(word) > 2) %>%
```

```
ggplot(aes(tf_idf, fct_reorder(word, tf_idf), fill = year)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~year, ncol = 2, scales = "free") +
  labs(x = "tf-idf", y = NULL)
```



So that gives an idea which words are frequent and unique to certain documents.

Now let's switch gears to quantda for some additional word relationship tools. We'll also get into some ways to assess the similarity of documents.

```
tokens <- tokens(epa_corp, remove_punct = TRUE)
toks1 <- tokens_select(tokens, min_nchar = 3)
toks1 <- tokens_tolower(toks1)
toks1 <- tokens_remove(toks1, pattern = (stop_vec))
dfm <- dfm(toks1)

#first the basic frequency stat
tstat_freq <- textstat_frequency(dfm, n = 5, groups = year)
head(tstat_freq, 10)
```

```
##           feature frequency rank docfreq group
## 1  environmental      127     1         1  2015
## 2   communities       99     2         1  2015
## 3             epa       92     3         1  2015
## 4         justice       84     4         1  2015
## 5      community       47     5         1  2015
```

```
## 6 environmental      109    1    1 2016
## 7 communities       85    2    1 2016
## 8 justice           71    3    1 2016
## 9 epa               48    4    1 2016
## 10 federal          31    5    1 2016
```

Another useful word relationship concept is that of the n-gram, which essentially means tokenizing at the multi-word level

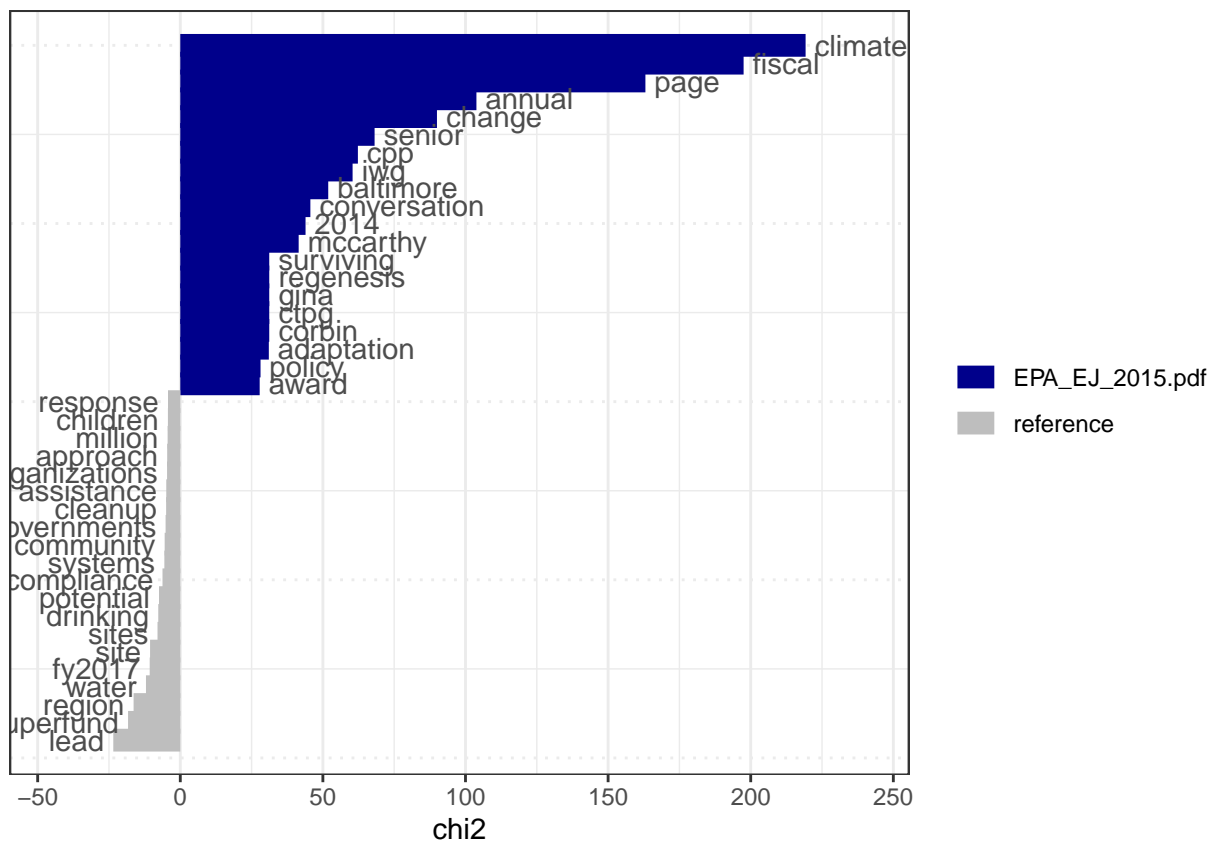
```
toks2 <- tokens_ngrams(toks1, n=2)
dfm2 <- dfm(toks2)
dfm2 <- dfm_remove(dfm2, pattern = c(stop_vec))
freq_words2 <- textstat_frequency(dfm2, n=20)
freq_words2$token <- rep("bigram", 20)
freq_words2
```

```
##           feature frequency rank docfreq group token
## 1 environmental_justice    556    1     6 all bigram
## 2 technical_assistance    139    2     6 all bigram
## 3 drinking_water         133    3     6 all bigram
## 4 public_health          123    4     6 all bigram
## 5 progress_report        108    5     6 all bigram
## 6 air_quality            73    6     6 all bigram
## 7 water_systems          66    7     6 all bigram
## 8 vulnerable_communities   65    8     6 all bigram
## 9 epa_region             62    9     5 all bigram
## 10 environmental_public    57   10     6 all bigram
## 11 federal_agencies       56   11     6 all bigram
## 12 national_environmental  51   12     6 all bigram
## 13 justice_fy2017         51   12     1 all bigram
## 14 fy2017_progress        51   12     1 all bigram
## 15 superfund_sites        48   15     4 all bigram
## 16 indigenous_peoples     46   16     6 all bigram
## 17 civil_rights           46   16     5 all bigram
## 18 local_governments      45   18     6 all bigram
## 19 urban_waters           44   19     6 all bigram
## 20 overburdened_communities 43   20     6 all bigram
```

```
#tokens1 <- tokens_select(tokens1, pattern = stopwords("en"), selection = "remove")
```

Now we can upgrade that by using all of the frequencies for each word in each document and calculating a chi-square to see which words occur significantly more or less within a particular target document

```
keyness <- textstat_keyness(dfm, target = 1)
textplot_keyness(keyness)
```




And finally, we can run a hierarchical clustering algorithm to assess document similarity. This tends to be more informative when you are dealing with a larger number of documents, but we'll add it here for future reference.

```
dist <- as.dist(textstat_dist(dfm))
clust <- hclust(dist)
plot(clust, xlab = "Distance", ylab = NULL)
```



## Cluster Dendrogram



Distance  
hclust (\*, "complete")

1. What are the most frequent trigrams in the dataset? How does this compare to the most frequent bigrams? Which n-gram seems more informative here, and why?

```
toks3 <- tokens_ngrams(toks1, n=3)
dfm3 <- dfm(toks3)
dfm3 <- dfm_remove(dfm3, pattern = c(stop_vec))
freq_words3 <- textstat_frequency(dfm3, n=20)
freq_words3$token <- rep("trigram", 20)
freq_words2
```

##	feature	frequency	rank	docfreq	group	token
## 1	environmental_justice	556	1	6	all	bigram
## 2	technical_assistance	139	2	6	all	bigram
## 3	drinking_water	133	3	6	all	bigram
## 4	public_health	123	4	6	all	bigram
## 5	progress_report	108	5	6	all	bigram
## 6	air_quality	73	6	6	all	bigram
## 7	water_systems	66	7	6	all	bigram
## 8	vulnerable_communities	65	8	6	all	bigram
## 9	epa_region	62	9	5	all	bigram
## 10	environmental_public	57	10	6	all	bigram
## 11	federal_agencies	56	11	6	all	bigram

## 12	national_environmental	51	12	6	all bigram
## 13	justice_fy2017	51	12	1	all bigram
## 14	fy2017_progress	51	12	1	all bigram
## 15	superfund_sites	48	15	4	all bigram
## 16	indigenous_peoples	46	16	6	all bigram
## 17	civil_rights	46	16	5	all bigram
## 18	local_governments	45	18	6	all bigram
## 19	urban_waters	44	19	6	all bigram
## 20	overburdened_communities	43	20	6	all bigram

```
freq_words3
```

##	feature	frequency	rank	docfreq	group	token
## 1	justice_fy2017_progress	51	1	1	all	trigram
## 2	fy2017_progress_report	51	1	1	all	trigram
## 3	environmental_public_health	50	3	6	all	trigram
## 4	environmental_justice_fy2017	50	3	1	all	trigram
## 5	national_environmental_justice	37	5	6	all	trigram
## 6	office_environmental_justice	32	6	6	all	trigram
## 7	epa's_environmental_justice	32	6	6	all	trigram
## 8	environmental_justice_progress	30	8	4	all	trigram
## 9	justice_progress_report	30	8	4	all	trigram
## 10	environmental_justice_concerns	30	8	5	all	trigram
## 11	drinking_water_systems	29	11	5	all	trigram
## 12	annual_environmental_justice	27	12	5	all	trigram
## 13	environmental_justice_advisory	27	12	6	all	trigram
## 14	fiscal_annual_environmental	25	14	3	all	trigram
## 15	justice_advisory_council	24	15	6	all	trigram
## 16	environmental_justice_grants	22	16	5	all	trigram
## 17	technical_assistance_communities	20	17	6	all	trigram
## 18	communities_environmental_justice	20	17	5	all	trigram
## 19	safe_drinking_water	19	19	5	all	trigram
## 20	technical_assistance_services	19	19	5	all	trigram

```
#tokens1 <- tokens_select(tokens1,pattern = stopwords("en"), selection = "remove")
```

The most popular tri-grams are justice\_fy2017\_progress, fy2017\_progress\_report and environmental\_public\_health. The most popular bi-grams are environmental\_justice, technical\_assistance and drinking\_water. In this case, the bi-grams seem much more useful because the tri-grams contain a lot of noise with most trigrams contains repeats of same keywords arranged differently.

**2. Choose a new focal term to replace “justice” and recreate the correlation table and network (see corr\_paragraphs and corr\_network chunks). Explore some of the plotting parameters in the cor\_network chunk to see if you can improve the clarity or amount of information your plot conveys. Make sure to use a different color for the ties!**

```
word_cors <- par_words %>%
  add_count(par_id) %>%
  filter(n >= 50) %>%
```

```

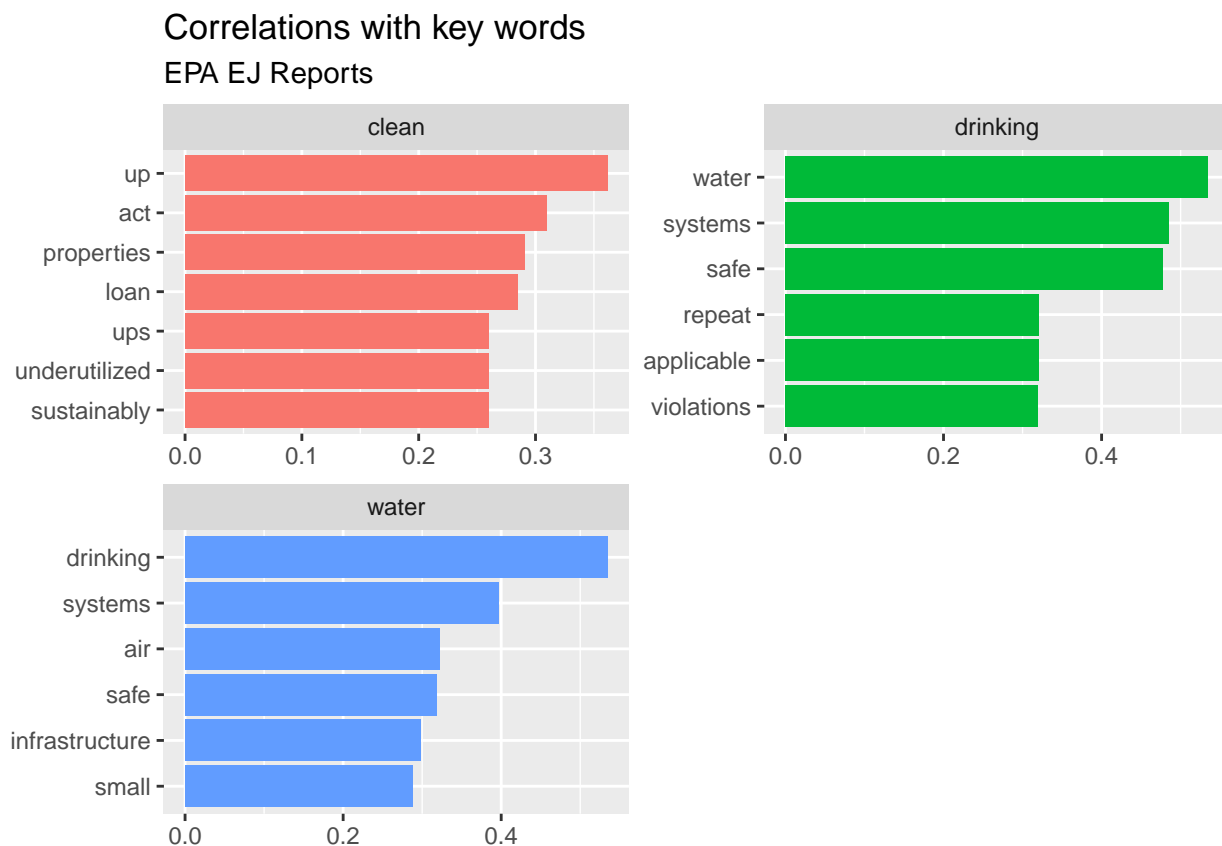
select(-n) %>%
pairwise_cor(word, par_id, sort = TRUE)

wat_cors <- word_cors %>%
  filter(item1 == "water")

word_cors %>%
  filter(item1 %in% c("clean", "water", "scarcity", "drinking"))%>%
  group_by(item1) %>%
  top_n(6) %>%
  ungroup() %>%
  mutate(item1 = as.factor(item1),
         name = reorder_within(item2, correlation, item1)) %>%
  ggplot(aes(y = name, x = correlation, fill = item1)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~item1, ncol = 2, scales = "free")+
  scale_y_reordered() +
  labs(y = NULL,
       x = NULL,
       title = "Correlations with key words",
       subtitle = "EPA EJ Reports")

```

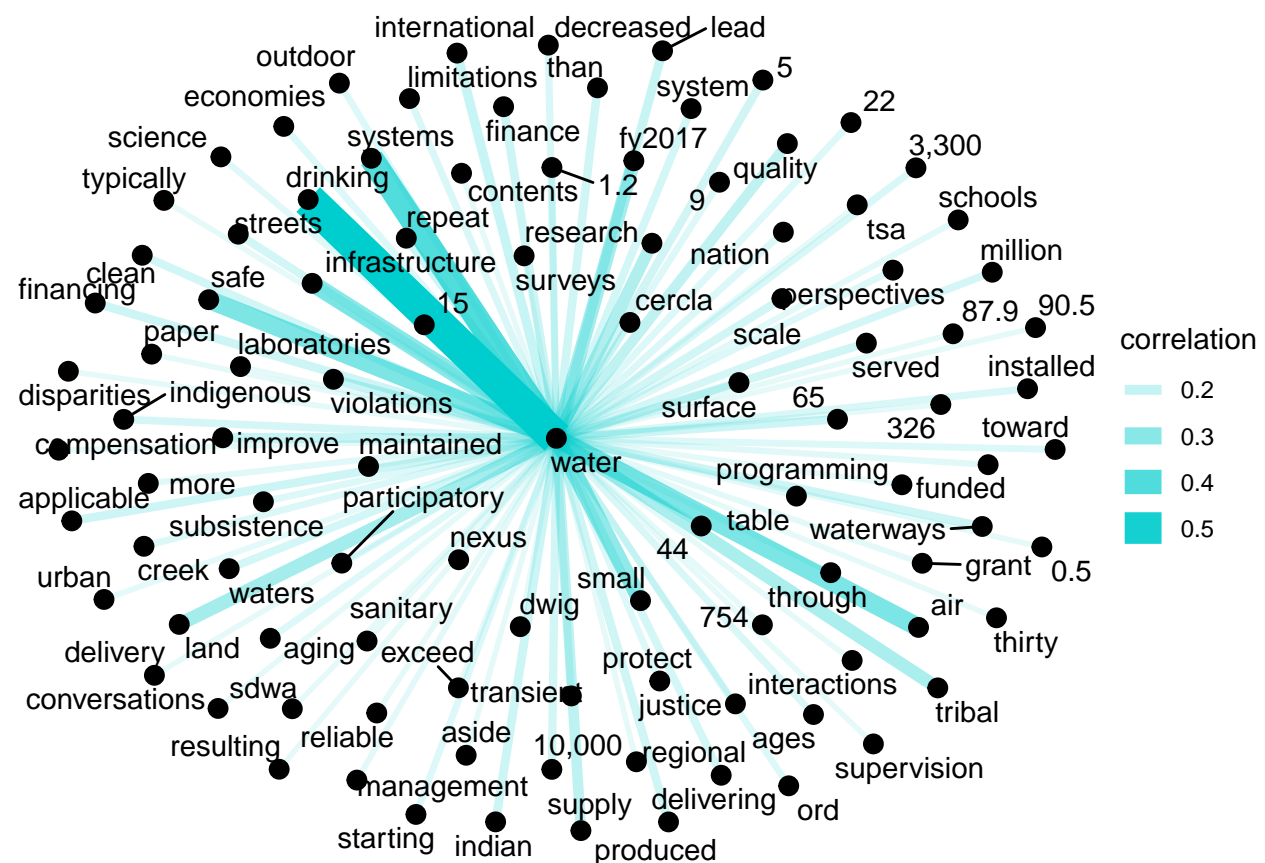
## Selecting by correlation



```
#let's zoom in on just one of our key terms
water_cors <- word_cors %>%
  filter(item1 == "water") %>%
  mutate(n = 1:n())
```

## Water as a focal term


```
water_cors %>%
  filter(n <= 100) %>%
  graph_from_data_frame() %>%
  ggraph(layout = "fr") +
  geom_edge_link(aes(edge_alpha = correlation, edge_width = correlation), edge_colour = "c") +
  geom_node_point(size = 3) +
  geom_node_text(aes(label = name), repel = TRUE,
    point.padding = unit(0.3, "lines")) +
  theme_void()
```




3. Write a function that allows you to conduct a keyness analysis to compare two individual EPA reports (hint: that means target and reference need to both be individual reports). Run the function on 3 pairs of reports, generating 3 keyness plots.

```
keyness_comp <- function(report1, report2){  
  files <- c(report1,report2)  
  reports <- lapply(files, pdf_text)  
  pdf <- readtext(files, docvarsfrom = "filenames",  
                  docvarnames = c("type", "subj", "year"),  
                  sep = "_")  
  
  #creating an initial corpus containing our data  
  corp <- corpus(x = pdf, text_field = "text" )  
  tokens <- tokens(corp, remove_punct = TRUE)  
  toks1<- tokens_select(tokens, min_nchar = 3)  
  toks1 <- tokens_tolower(toks1)  
  toks1 <- tokens_remove(toks1, pattern = (stop_vec))  
  dfm <- dfm(toks1)  
  keyness <- textstat_keyness(dfm, target = 1)  
  
  return(keyness)  
}
```


```
keyness_computed_1<-keyness_comp(here("dat", "EPA_EJ_2016.pdf"), here("dat", "EPA_EJ_2018.pdf"))  
textplot_keyness(keyness_computed_1)
```



```
keyness_computed_2<-keyness_comp(here("dat", "EPA_EJ_2016.pdf"), here("dat", "EPA_EJ_2019.pdf"))
textplot_keyness(keyness_computed_2)
```



```
keyness_computed_3<-keyness_comp(here("dat", "EPA_EJ_2019.pdf"), here("dat", "EPA_EJ_2020.pdf"))
textplot_keyness(keyness_computed_3)
```



4. Select a word or multi-word term of interest and identify words related to it using windowing and keyness comparison. To do this you will create two objects: one containing all words occurring within a 10-word window of your term of interest, and the second object containing all other words. Then run a keyness comparison on these objects. Which one is the target, and which the reference?

Water Quality as multi-word term of interest

```
water_quality <- c("water", "water quality")
toks_inside <- tokens_keep(tokens, pattern = water_quality, window = 10)
toks_inside <- tokens_remove(toks_inside, pattern = water_quality) # remove the keywords
toks_outside <- tokens_remove(tokens, pattern = water_quality, window = 10)
```

```
dfmat_inside <- dfm(toks_inside)
dfmat_outside <- dfm(toks_outside)

tstat_key_inside <- textstat_keyness(rbind(dfmat_inside, dfmat_outside),
                                     target = seq_len(ndoc(dfmat_inside)))
head(tstat_key_inside, 25)
```

##	feature	chi2	p	n_target	n_reference
## 1	drinking	1983.86441	0.000000e+00	138	3



## 2	systems	1026.45693	0.000000e+00	92	25
## 3	infrastructure	197.22605	0.000000e+00	36	45
## 4	small	192.98029	0.000000e+00	43	69
## 5	system	182.53698	0.000000e+00	27	23
## 6	safe	178.65850	0.000000e+00	24	17
## 7	health-based	161.52831	0.000000e+00	13	1
## 8	land	115.86858	0.000000e+00	23	30
## 9	clean	113.60064	0.000000e+00	30	58
## 10	served	75.98434	0.000000e+00	9	4
## 11	tribal	75.41803	0.000000e+00	49	198
## 12	applicable	73.25277	0.000000e+00	6	0
## 13	quality	60.87612	6.106227e-15	33	118
## 14	air	54.92440	1.252332e-13	44	202
## 15	supply	52.24150	4.908296e-13	6	2
## 16	supervision	47.42411	5.717538e-12	5	1
## 17	finance	45.26037	1.725042e-11	6	3
## 18	non-community	44.17991	2.995404e-11	4	0
## 19	utilities	39.68904	2.977918e-10	6	4
## 20	surface	39.40271	3.448182e-10	5	2
## 21	delivering	37.70655	8.222780e-10	7	7
## 22	violations	37.70655	8.222780e-10	7	7
## 23	financing	33.97142	5.592754e-09	4	1
## 24	repeat	33.40363	7.488399e-09	5	3
## 25	wastewater	32.04017	1.510176e-08	9	16

The tokens inside the 10-word window are the target while the tokens outside the 10-word window are the reference.